

CS M213A / ECE M202A (Fall 2025)

From Words to Shields: Agentic Privacy for Smart Sensors

Daniel Luzzatto, Sam Seban

Project Mentor: Brian Wang

Outline

- Motivation / Problem Statement
- Related / Prior Work
- Our Goals & Contributions
- Team Contributions
- Technical Method / Algorithm
- Results & Findings
- Conclusions & Future Work
- Live Demo

Motivation

- Smart environments stream raw audio/video to cloud services → privacy risk.
- Users cannot express simple privacy rules (e.g., “blur faces,” “mute medical terms”).
- Existing solutions rely on manual, rigid pipelines that don’t scale or adapt to changing policies or are not robust.
- Need a system that enforces privacy automatically from natural-language preferences.

Project Goals

- Translate user natural-language privacy requirements into executable pipelines.
- Use an LLM-based agent to build and run privacy operators for audio, video.
- Ensure consistent behavior across modalities and live/non-live inputs.
- Provide safety, verification, and auditability.

Related / Prior Work

Peekaboo (Jin et al., 2022)

- Enforces privacy using a developer-written manifest and a fixed set of detectors/transforms.
- Limitation: Static and inflexible, it cannot create new logic or tools.
- Relevance to our work: We keep the pipeline idea (detect → transform → verify) but extend it with dynamic tool generation and LLM-based planning, eliminating the static manifest bottleneck.

1-2-3 Check (Li et al., 2024)


- Multi-agent system that enforces Contextual Integrity (CI) norms in text/audio.
- Uses a Checker Agent to validate outputs and reduce leakage.
- Limitation: Focuses only on contextual leakage in text/audio; no multimodal sanitization, no tool generation.
- Relevance to our work: Inspires our policy verification layer, but our system goes beyond CI checking by supporting vision + audio and generating executable sanitization tools on demand.

Our Novelty

- Autonomous tool generation when existing operators aren't enough.
- Unified policy framework covering audio, video, and multimodal inputs.
- Dynamic verification with automatic retries and full replanning on failure.
- Full auditability via manifests and provenance logs.
- Combines LLM flexibility with strict verification, making it adaptable and trustworthy.

High-Level Approach

- Ingest raw sensor data (video/audio) and user's natural-language privacy request
- Planner Agent interprets the request and produces a manifest (execution plan)
- Checks tool registry; reuses existing tools or calls Tool Generator to create new ones
- Executor Agent runs the tools in sequence to transform/sanitize the data
- Verification Module checks the output against the privacy policy
- All steps logged for auditability



Team Contributions

Sam

- Speech processing pipeline
- Tool generation module
- Metrics module
- Audit & logging framework
- Executor runtime

Daniel


- LLM planner agent
- Face detection & blurring pipeline
- Live-streaming redaction
- Verification modules (video + audio)
- Closed-loop recovery design

Collaborative Work

- Weekly meetings
- Joint code reviews
- Shared testing & integration sessions

Algorithm - Closed-loop

- After execution, output goes through the Verification Module
- If verification passes: output is released + logged
- If verification fails: Automatic retries using fallback tools or alternative configurations; If still failing → system triggers full replanning
- Planner receives diagnostic report and updates the manifest or generates new tools
- Ensures privacy requirements are always met before data leaves the system



Algorithm - Planning


- Uses llama-3.3-70b-versatile to convert natural-language privacy requests into JSON manifests
- Chooses tools from registry (face blur, keyword mute, etc.) or proposes new tools via naming templates
- Manifest encodes a pipeline of steps with tool names + arguments
- System checks tool availability; missing tools → auto generation
- Final manifest includes metadata on generated tools and is sent to execution

User Request: “Blur faces and remove license plates”

```
{  
    "pipeline": [  
        {"tool": "blur_faces", "args": {}},  
        {"tool": "blur_license_plates", "args": {}}  
    ]  
}
```


Algorithm - Face blurring

- Hybrid detection pipeline: YuNet, KCF tracker, Kalman Filter
- Cascaded strategy: detect every N frames + track in between
- Kalman model: constant velocity, 8-state filter, ~60-frame prediction cap
- Grayscale preprocessing: CLAHE + sharpening for higher contrast
- Live vs Offline:
 - Live → single-pass
 - Offline → two-pass (detect first, verify, then blur)
- Verification (offline):
 - Miss-ratio threshold (<10%)
 - Laplacian-variance blur strength (<50)



Algorithm - Speech beeping

- Whisper (base) → transcript + timestamps
- llama-3.3-70b extracts sensitive phrases from transcript based on user intent
- Sliding-window match on timestamped words with precise time intervals
- pydub mutes or inserts a beep in detected intervals
- Live streaming: FasterWhisperASR + small buffer, Real-time keyword detection before playback.
- Verification (offline):
 - Temporal integrity check (valid timestamps, no long segments, low overlap)
 - Compliance check: re-run ASR → confirm sensitive phrases no longer appear



Algorithm - Tool Generation



- Triggered when the Planner requests a tool not found in the registry
- Uses LLM with a specialized system prompt containing: allowed libraries whitelist, code template for a valid tool, guidelines for video/audio processing patterns. processing patterns
- LLM outputs full tool Python code
- System cleans the code and saves it under the correct category
- Tool is dynamically imported into the registry (no restart required)
- Newly generated tool becomes immediately available for execution in the pipeline



Evaluation Metrics


Category	Metrics
Correctness	Face-blur accuracy; mute/beep correctness
Robustness	Pipeline success rate; tool-generation success %; recovery success
Adaptability	Time to generate new pipeline; # of LLM steps; success rate on unseen requirements
Performance	End-to-end latency; resource usage (CPU/GPU); throughput
Security	Sandbox compliance; manifest verification; generated code safety
Auditability	Log completeness; manifest accuracy; decision traceability
Baseline Comparison	Latency overhead caused by LLM calls

Results



Results

Comparison of time spent in each stage between rigid traditional pipelines and adaptive LLM-based approach.



Analysis & Discussion of Results

- Overall success rate before any retry attempts is **59%**.
- When the recovery loop is triggered, an average recovery success rate of **24%** is observed. As expected, all three success rates systematically decrease as scenario complexity increases.
- In the most complex case ("Vid 3"), the `mute_keywords` tool struggled the most and frequently failed verification due to overlapping speech making automatic transcription more challenging.
- ~10.2s overhead compared to traditional pipelines
- More latency than traditional pipelines due to LLM calls, error recovery mechanism

Conclusions

- Built an end-to-end system where LLM agents translate natural-language privacy policies into executable pipelines.
- Achieved 69% overall success, with the retry loop recovering 24% of failed attempts.
- Strong multimodal performance:
 - Face tracking with <10% miss ratio.
 - Keyword muting/redaction across diverse privacy intents.
- Cuts implementation time from weeks of engineering → seconds of automated pipeline synthesis.

Future Work & Next Directions

Performance & Deployment

- GPU acceleration for YuNet + Whisper to reduce end-to-end latency
- Exploration of smaller/distilled models (e.g., Whisper-tiny + fine-tuning) for on-device inference
- Improved handling of multi-speaker audio via speaker diarization

Reliability & Verification

- Fine-tuning a dedicated model for privacy-tool code generation
- Human-in-the-loop verification (e.g., MTurk) to train stronger automated verifiers

Scope Expansion

- Extend beyond audio–video to arbitrary sensor modalities

References

1. Jin, H., et al. "Peekaboo: A Hub-Based Approach to Enable Transparency in Data Processing within Smart Homes." [arXiv:2204.04540](https://arxiv.org/abs/2204.04540)
2. Li, W., et al. "1-2-3 Check: Contextual Integrity-Aware Multi-Agent Privacy Reasoning." [arXiv:2508.07667](https://arxiv.org/abs/2508.07667)
3. Iravantchi, Y., et al. "PrivacyLens: On-Device PII Removal from RGB Images Using Thermal Sensing." [PETS 2024](https://www.pets-conf.org/pets2024/)
4. Zhou, J., et al. "Privacy-sensitive Objects Pixelation for Live Video Streaming." [arXiv:2101.00604](https://arxiv.org/abs/2101.00604)
5. Zhang, S., et al. "Evaluating the Efficacy of Large Language Models for Generating Fine-Grained Visual Privacy Policies in Homes." [arXiv:2508.00321](https://arxiv.org/abs/2508.00321)

Live Demo

Thank You