CS M213A / ECE M202A (Fall 2025)

From Words to Shields: Agentic Privacy for Smart Sensors

Sam E. Seban & Daniel Luzzatto

Motivation and Objectives

Project Outline:

- Smart devices constantly collect audio/video, often sending raw data to the cloud.
- Users can't control what is filtered out or how privacy is enforced.
- We want a system where users can simply *say what they want protected* ("blur children's faces", "mute medical terms"), and the system automatically enforces it.

Audience:

- Protects people in homes, hospitals, workplaces…
- Helps meet changing regulations (e.g., GDPR).
- Reduces engineering effort for companies.

Goals:

- A prototype "agentic privacy hub" where an LLM builds and executes privacy pipelines.
- Multimodal support: audio, video, text, tabular.
- Verification system ensuring sanitized data actually respects the request.
- Quantitative evaluation of correctness, robustness, latency, and adaptability.

Technical Approach and Novelty

How it's done today:

Systems like Peekaboo use fixed, developer-built pipelines. They provide strong guarantees, but low flexibility, as new privacy rules require new code.

Our high-level approach:

- LLM agent interprets natural-language privacy rules.
- It builds a pipeline: selecting tools, generating missing ones, executing, and verifying.
- Closed-loop: replans if verification fails.

Novelty:

- Autonomous tool generation (sandboxed).
- Unified policy for video, audio, text, tabular.
- Dynamic verification + recovery (retry → replan).
- Strong auditability (manifests, logs).

Methods

- Vision: face detection and blur using OpenCV, YuNet, Kalman Filter.
- Speech: keywords detection and beep/silence using Whisper.
- LLM for pipeline and tool generation: llama-3.3-70b-versatile.

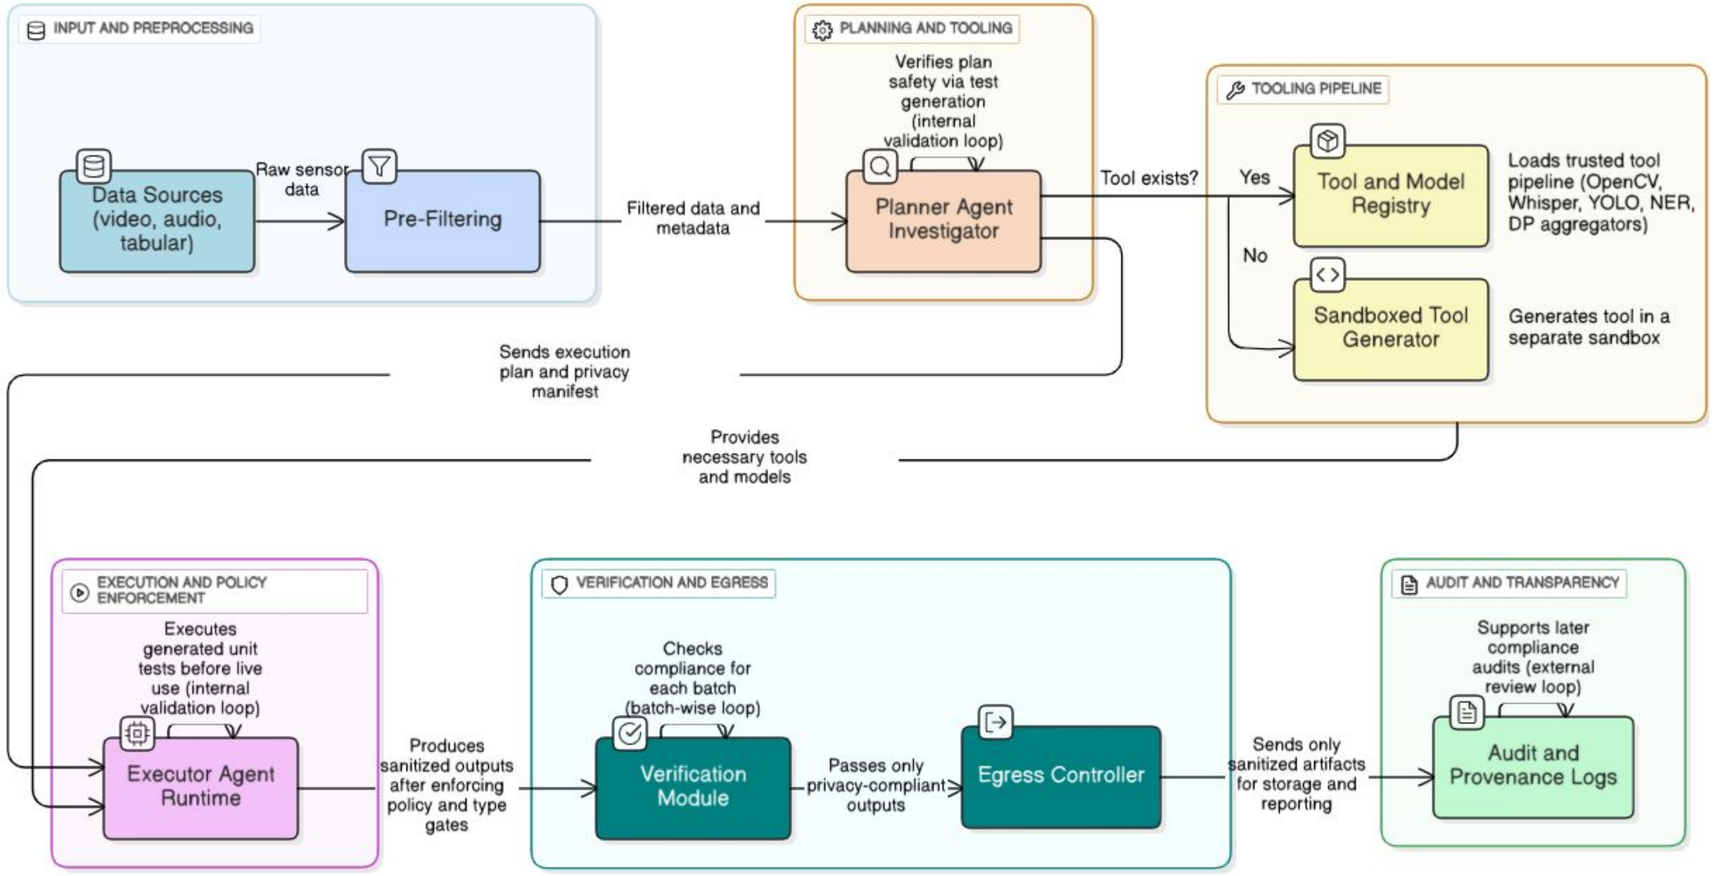

Evaluation and Metrics

Category	Metrics
Correctness	Face-blur accuracy; keyword F1; mute/beep correctness
Robustness	Pipeline success rate; tool-generation success %; recovery success
Adaptability	Time to generate new pipeline; # of LLM steps
Performance	End-to-end latency; resource usage
Security	Sandbox compliance; manifest verification

Current Status and Next Steps

Current status:

- Implemented face recognition and blur tool and verification for non-live videos.
- Implemented keyword detection and beep/silence and verification for non-live audio files.
- Implemented pipeline planner.
- Implemented tool generator.

Next steps:

- Add multimodal, and tabular input tools.
- Adapt the existing tools for live applications.
- Implement verification of the generated tools.
- Implement the closed-loop architecture.
- Run the generated code in a sandbox environment for security.

Prompt: “Blur faces”

Generated manifest:

```
{
  "pipeline": [
    {
      "tool": "blur_faces",
      "args": {
        "kernel": 31
      }
    }
  ]
}
```

Prompt: “Hide the background”

Generated manifest:

```
{
  "pipeline": [
    {
      "tool": "remove_background",
      "args": {
        "background_option": "blur"
      }
    }
  ]
}
```

Tool doesn’t exist

↓

Trigger tool generation